

Assignment 4 A Barrier Mechanism in Linux Kernel (150 points)

Exercise Objectives

1. To learn the basic programming technique in Linux kernel.
2. To learn the software structures of kernel thread and synchronization.
3. To practice the mechanisms and data structures of system call, wait queue, and other Linux kernel objects

Project Assignment

Barrier synchronization has been applied widely in parallel computing to synchronize the execution of parallel loops. The basic idea is to allow parallel loop computation to start next iteration if all loops have completed the previous iteration. In NPTL pthread library, it is realized by 3 functions for threads of the same process, i.e., `pthread_barrier_init`, `pthread_barrier_wait`, and `pthread_barrier_destroy`. A description of the functions can be found in <https://docs.oracle.com/cd/E19253-01/816-5137/gfwek/index.html>, and their source code is available in glibc NPTL cross reference page <http://code.metager.de/source/xref/gnu/glibc/nptl/>.

In this assignment, you are required to develop a barrier synchronization mechanism in Linux kernel which can be invoked by user processes via system call interface. The mechanism consists of 3 Linux kernel system calls which are resemble to the pthread barrier:

1. `barrier_init(unsigned int count, unsigned int *barrier_id, signed int timeout)` – to initiate a barrier in the caller's address space. "count" is the number of threads (i.e., processes and lightweight processes) that are required to be synchronized. If the barrier is initiated successfully, an integer barrier id is returned to the caller process. The `timeout` parameter indicates the maximal waiting time in nanoseconds of any participating threads. A zero or negative value in `timeout` implies there is no timeout in the barrier.
2. `barrier_wait(unsigned int barrier_id)` – to wait for a barrier synchronization. The caller process will be blocked until the required number of threads have called the wait function, or a timeout occurs, i.e., the barrier synchronization fails.
3. `barrier_destroy(unsigned int barrier_id)` – to destroyed the barrier of the id `barrier_id`.

Your implementation should meet the following requirements:

- The calls should return proper error status if the required operations cannot be completed. For instance, when a timeout occurs, the synchronization fails. All waiting threads should be waked up and return with an error status.
- Each barrier should be initiated in caller's address space and can only be used by the processes sharing the space.
- `barrier_id` is unique in its address space and the `barrier_id`'s of different address space may have the same value.
- Once a barrier is initiated, it can be used for consecutive synchronization operations until it is destroyed.
- The barrier mechanism must work properly in SMP systems and allow multiple user processes work on multiple barriers in their corresponding address space concurrently. Unfortunately, Quark doesn't have multiple cores. Thus, we cannot test this execution environment.

To demonstrate your barrier implementation, a testing program should be developed. The testing program forks two child processes. In each child process, there are 5 threads to exercise the 1st barrier and additional 20 threads use the 2nd barrier. Each thread sleeps a random amount of time before

entering a new round of synchronization. The average sleep time in microsecond should be an input integer to the testing program. The testing program terminates after 100 rounds of synchronization on each barrier in each child process. To wait for the completion of child processes, your testing program can use “wait” system call. An example can be found in <http://www.csl.mtu.edu/cs4411.ck/www/NOTES/process/fork/wait.html>.

Since Linux system calls are statically defined and compiled into the kernel, you will need to rebuild the kernel with new system calls. You may consider the following suggested procedure to rebuild 3.19 Linux kernel for Galileo Gen2 board:

- Use the pre-built toolchain and the Linux kernel source linux3-19-r0.tar.bz2 from Dropbox/CSE438-530-Gen2_2017/Boot_image&SDK. Also, the “.config” file in the source code can be used for kernel building.
- Include the cross-compilation tools in your PATH:
export PATH=path_to_sdk/sysroots/x86_64-pokysdk-linux/usr/bin/i586-poky-linux:\$PATH
- Cross-compile the kernel
ARCH=x86 LOCALVERSION= CROSS_COMPILE=i586-poky-linux- make -j4
- Build and extract the kernel modules from the build to a target directory (e.g ../galileo-install)
ARCH=x86 LOCALVERSION= INSTALL_MOD_PATH=../galileo-install CROSS_COMPILE=i586-poky-linux- make modules_install
- Extract the kernel image (bzImage) from the build to a target directory (e.g ../galileo-install)
cp arch/x86/boot/bzImage ../galileo-install/
- Install the new kernel and modules from the target directory (e.g ../galileo-install) to your micro SD card
 - Replace the bzImage found in the first partition (ESP) of your micro SD card with the one from your target directory (backup the bzImage on the micro SD card e.g. rename it to bzImage.old)
 - Copy the kernel modules from the target directory to the /lib/modules/ directory found in the second partition of your micro SD card.
- Reboot into your new kernel

Note that, after the first build, there is no need to rebuild/replace the kernel modules if you don't make any changes in the loadable modules. Also, to avoid building the kernel for each code revision, you may implement the functions in a loadable driver module and use device file operations to invoke them. Once the functions are fully developed, you can then add them as system calls and rebuild the kernel.

Due Date

The due date for both parts 1 is 11:59pm, Nov. 28.

What to Turn in for Grading

- The submission includes a **patch file** that consists of changes to the original kernel source code, a testing program, a Makefile (to make the testing program), and a README file. Compress the files into a zip archive named **cse530-teamX-assgn04.zip**. Note that any object code, original Linux kernel source code, or temporary build files should not be included in the submission. A submission that contains the original Linux kernel source code will not be graded and will receive a zero score. Please submit the zip archive to Blackboard by the due date and time.

- Please make sure that you comment the source files properly and the readme file includes a description about how to make and use your software. **Don't forget to add each team member's name and ASU id in the readme file.**
- There will be 20 points penalty per day if the submission is late. Note that submissions are time stamped by Blackboard. **If you have multiple submissions, only the newest one will be graded.** If needed, you can send an email to the instructor and TA to drop a submission.
- Your team must work on the assignment without any help from other teams and is responsible to the submission in Blackboard. **No collaboration between teams is allowed, except the open discussion in the forum on Blackboard.**
- Failure to follow these instructions may cause deduction of points.
- Here are few general rule for deductions:
 - No make file or compilation error -- 0 point for the part of the assignment.
 - Must have "-Wall" flag for compilation -- 5-point deduction for each warning.
 - 10-point deduction if no compilation or execution instruction in README file.
 - Source programs are not commented properly -- 10-20-point deduction.
- ASU Academic Integrity Policy (<http://provost.asu.edu/academicintegrity>), and FSE Honor Code (<http://engineering.asu.edu/integrity>) are strictly enforced and followed. **A grade XE will be assigned to any cases of AIP violation.**